



X86 memory

coherency

Background/Motivation



JVM
(Java Virtual Machine)

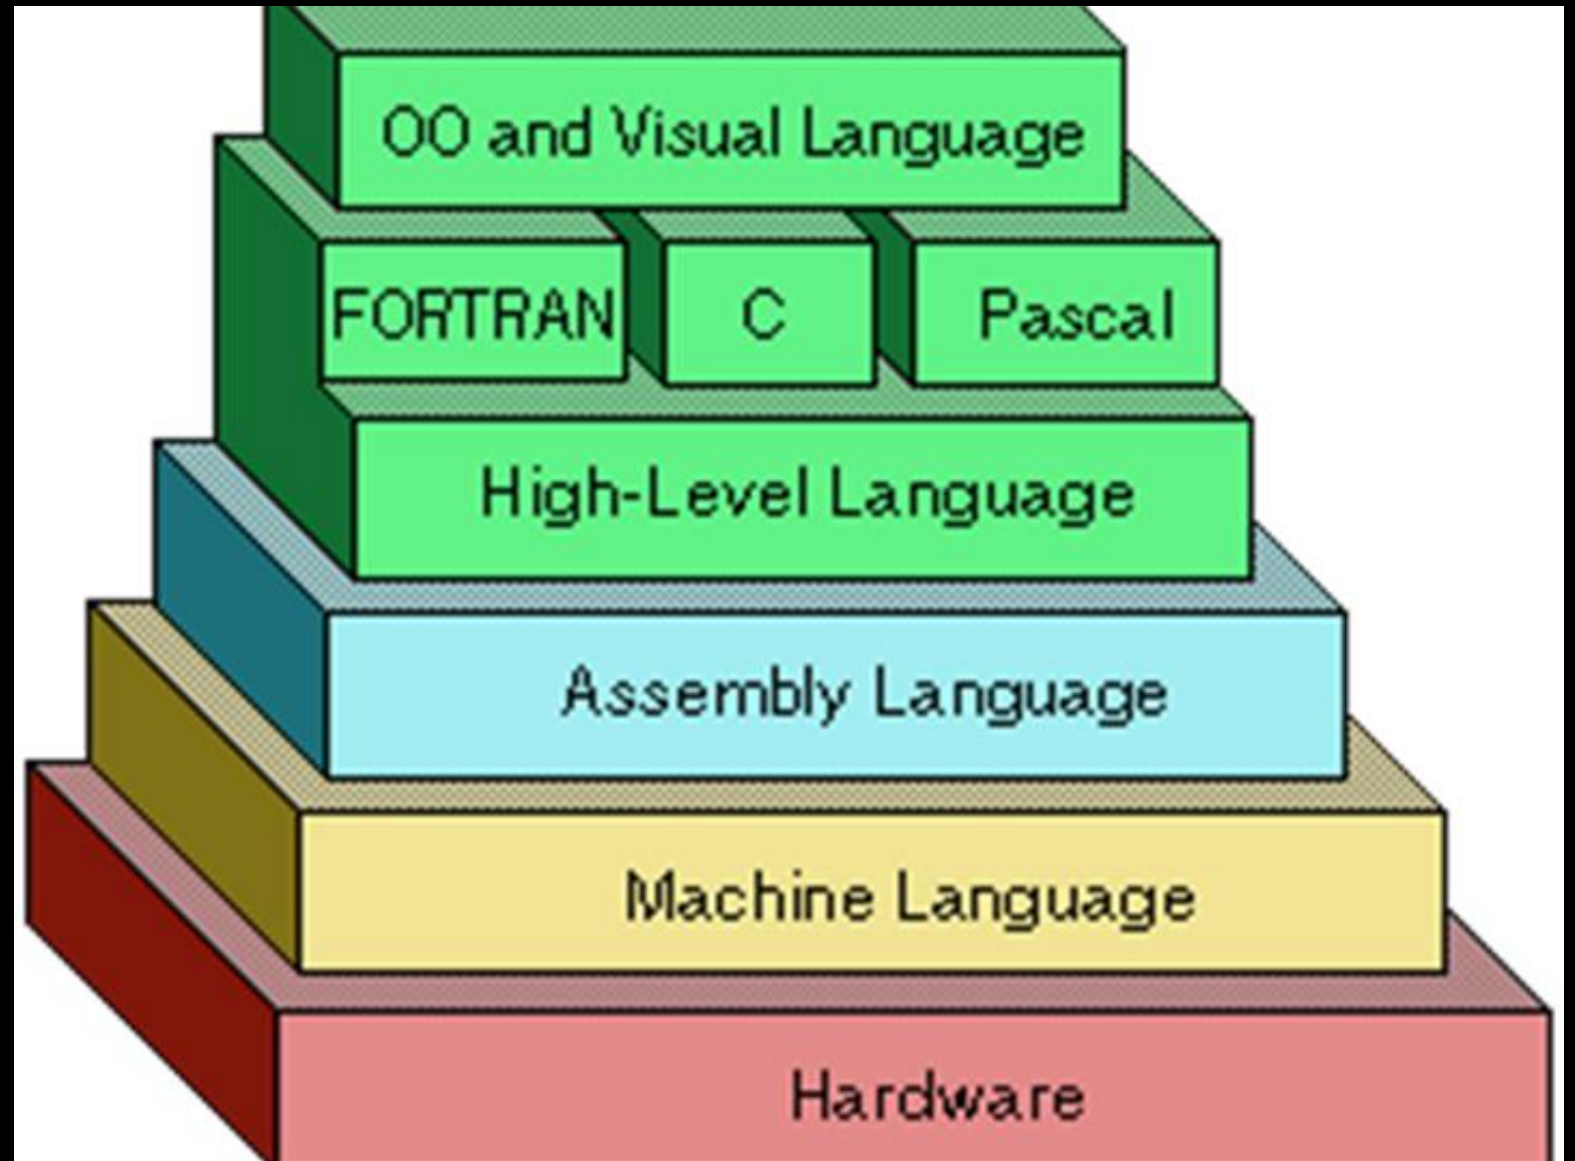
*Virtual machines are just C/C++ programs running on real Machines
like X86*




JVM
(Java Virtual Machine)

Understand from bottom up perspective (Why)


- Initially out of interest and curiosity
- But turned out to be useful in day job too



Multiprocessors
reorder memory
operations in
unintuitive ways



This behaviour is
necessary for
performance but
affects correctness
too!



We only need to care
about what we can
observe in software

Learning through experiments

Initially $A = B = 0$

Thread 1

`A = 1`

```
if (B == 0)
    print "Hello";
```

Thread 2

`B = 1`

```
if (A == 0)
    print "World";
```

What can be printed?

Initially A = B = 0

Thread 1

```
A = 1
if (B == 0)
    print "Hello";
```

Thread 2

```
B = 1
if (A == 0)
    print "World";
```

What can be printed?

*Single writer for each location in
memory
Single reader for different location*

can it print hello world ?

Initially $A = B = 0$

Thread 1

$A = 1$

if ($B == 0$)

 print "Hello";

Thread 2

$B = 1$

if ($A == 0$)

 print "World";

What can be printed?

For line 2 to print 0

Print B must happen before $B = 1$ on Thread-2

Thread 1

(1) $A = 1$

(2) print(B)

Thread 2

(3) $B = 1$

(4) print(A)



For line 4 to print 0

Print A must happen before $A = 1$ on Threa-1

Initially $A = B = 0$

Thread 1

$A = 1$

if ($B == 0$)

 print "Hello";

Thread 2

$B = 1$

if ($A == 0$)

 print "World";

What can be printed?

For line 2 to print 0

Print B must happen before $B = 1$

Thread 1


(1) $A = 1$

(2) print(B)

Thread 2

(3) $B = 1$

(4) print(A)



For line 4 to print 0

Print A must happen before $A = 1$

```
// x and y are initialised to 0
```

CORE-0	CORE-1
--------	--------

x = 1	y = 1
-------	-------

r0 = y	r1 = x
--------	--------

1. Thread0 runs first all way through r0=0, r1=1

2. Thread1 runs first all way through r0=1, r1=0

3. any other interleaving operation can only result in r0=1 , r1=1

Interleaving possibilities

STEP-1	STEP-2	STEP-3	STEP-4	(R0, R1)
core-1: y=1	core-0: x=1	core-1: r1=x	core-0: r0=y	(1,1)
core-1: y=1	core-0: x=1	core-0: r1=y	core-1: r0=x	(1,1)
core-0: x=1	core-1: y=1	core-0: r0=y	core-1: r1=x	(1,1)
core-0: x=1	core-1: y=1	core-1: r1=x	core-0: r0=y	(1,1)

- How can r0 and r1 == 0
- A re-order happen

Who is Re-Ordering

Thread 1

(1) A = 1
(2) print(B)

Thread 2

(3) B = 1
(4) print(A)



Is it
compiler ?

Is it
hardware ?

Or Both

- C is simple and easier to reason
- So good starting point

```
// x and y are initialised to 0
```

CORE-0	CORE-1
x = 1	y = 1
r0 = y	r1 = x

- Should only terminate if r0 and r1 equal to zero
- A re-order happen

```

1
2  #include <pthread.h>
3  #include <stdio.h>
4
5  int x, y = 0;
6  int r0, r1;
7
8  void *core0 (void *arg)
9  {
10     x = 1;
11     r1 = y;
12     return 0;
13 }
14
15 void *core1 (void *arg)
16 {
17     y = 1;
18     r0 = x;
19     return 0;
20 }
21
22
23 int main (void)
24 {
25     pthread_t thread0, thread1;
26     while (1) {
27         x = y = 0;
28         //Start threads
29         pthread_create (&thread0, NULL, core0, NULL);
30         pthread_create (&thread1, NULL, core1, NULL);
31
32         //wait for threads to complete
33         pthread_join (thread0, NULL);
34         pthread_join (thread1, NULL);
35
36         if (r0 == 0 && r1 == 0) {
37             printf("(r0=%d, r1=%d)\n", r0, r1);
38             break;
39         }
40
41     }
42     return 0;
43 }

```

- Lets rule out compiler Re-ordering
- Small hack to stop GCC compiler from re-ordering

```

1
2  #include <pthread.h>
3  #include <stdio.h>
4
5  int x, y = 0;
6  int r0, r1;
7
8  void *core0 (void *arg)
9  {
10     x = 1;
11     asm volatile ("" ::: "memory"); // ensure GCC compiler will not reorder
12     r1 = y;
13     return 0;
14 }
15
16 void *core1 (void *arg)
17 {
18     y = 1;
19     asm volatile ("" ::: "memory"); // ensure GCC compiler will not reorder
20     r0 = x;
21     return 0;
22 }
23
24
25 int main (void)
26 {
27     pthread_t thread0, thread1;
28     while (1) {
29         x = y = 0;
30         //Start threads
31         pthread_create (&thread0, NULL, core0, NULL);
32         pthread_create (&thread1, NULL, core1, NULL);
33
34         //wait for threads to complete
35         pthread_join (thread0, NULL);
36         pthread_join (thread1, NULL);
37
38         if (r0 == 0 && r1 == 0) {
39             printf("(r0=%d, r1=%d)\n", r0, r1);
40             break;
41         }
42     }
43     return 0;
44 }
45

```

- Program still terminates
- Meaning a Re-order happen !

Looking at one level lower just to be sure

```
8 void *core0 (void *arg)
9 {
10     x = 1;
11     asm volatile ("" ::: "memory");
12     r1 = y;
13     return 0;
14 }
15
16 void *core1 (void *arg)
17 {
18     y = 1;
19     asm volatile ("" ::: "memory");
20     r0 = x;
21     return 0;
22 }
23
```

```
1 core0:
2     push    rbp
3     mov     rbp, rsp
4     mov     QWORD PTR [rbp-8], rdi
5     mov     DWORD PTR x[rip], 1    ## write value 1 to x
6     mov     eax, DWORD PTR y[rip] ## read contents of y into register
7     mov     DWORD PTR r1[rip], eax
8     mov     eax, 0
9     pop     rbp
10    ret
11 core1:
12    push    rbp
13    mov     rbp, rsp
14    mov     QWORD PTR [rbp-8], rdi
15    mov     DWORD PTR y[rip], 1    ## write value 1 to y
16    mov     eax, DWORD PTR x[rip] ## read contents of x into register
17    mov     DWORD PTR r0[rip], eax
18    mov     eax, 0
19    pop     rbp
20    ret
```


Has to be at hardware level

1. Writes to memory even slower than reads as more book keeping required to keep caches coherent

Has to be at hardware level

1. Remember this from last week memory is super slow
2. Writes to memory even slower as more book keeping required to keep caches coherent

Cache and memory access latency (based on an Intel i7 8700)



Coherent Memory

Coherent memory (L1/L2/L3 caches and RAM)

- **Caches in CPU cores are coherent so if one CPU writes to its L1 cache, x86 hardware guarantees other CPUs will observe the change too. As software devs we don't need to worry about coherency here**


Non Coherent Memory

Non coherent memory (Store buffers, Load buffers, TLB, maybe more)

- **Store buffers (visible effect and we need to handle this)**
- Load Buffer (no visible effect on software on x86)
- TLB (visible effect especially in mem mapped files (another presentation maybe))
- Maybe others I don't know yet but they don't affect our software

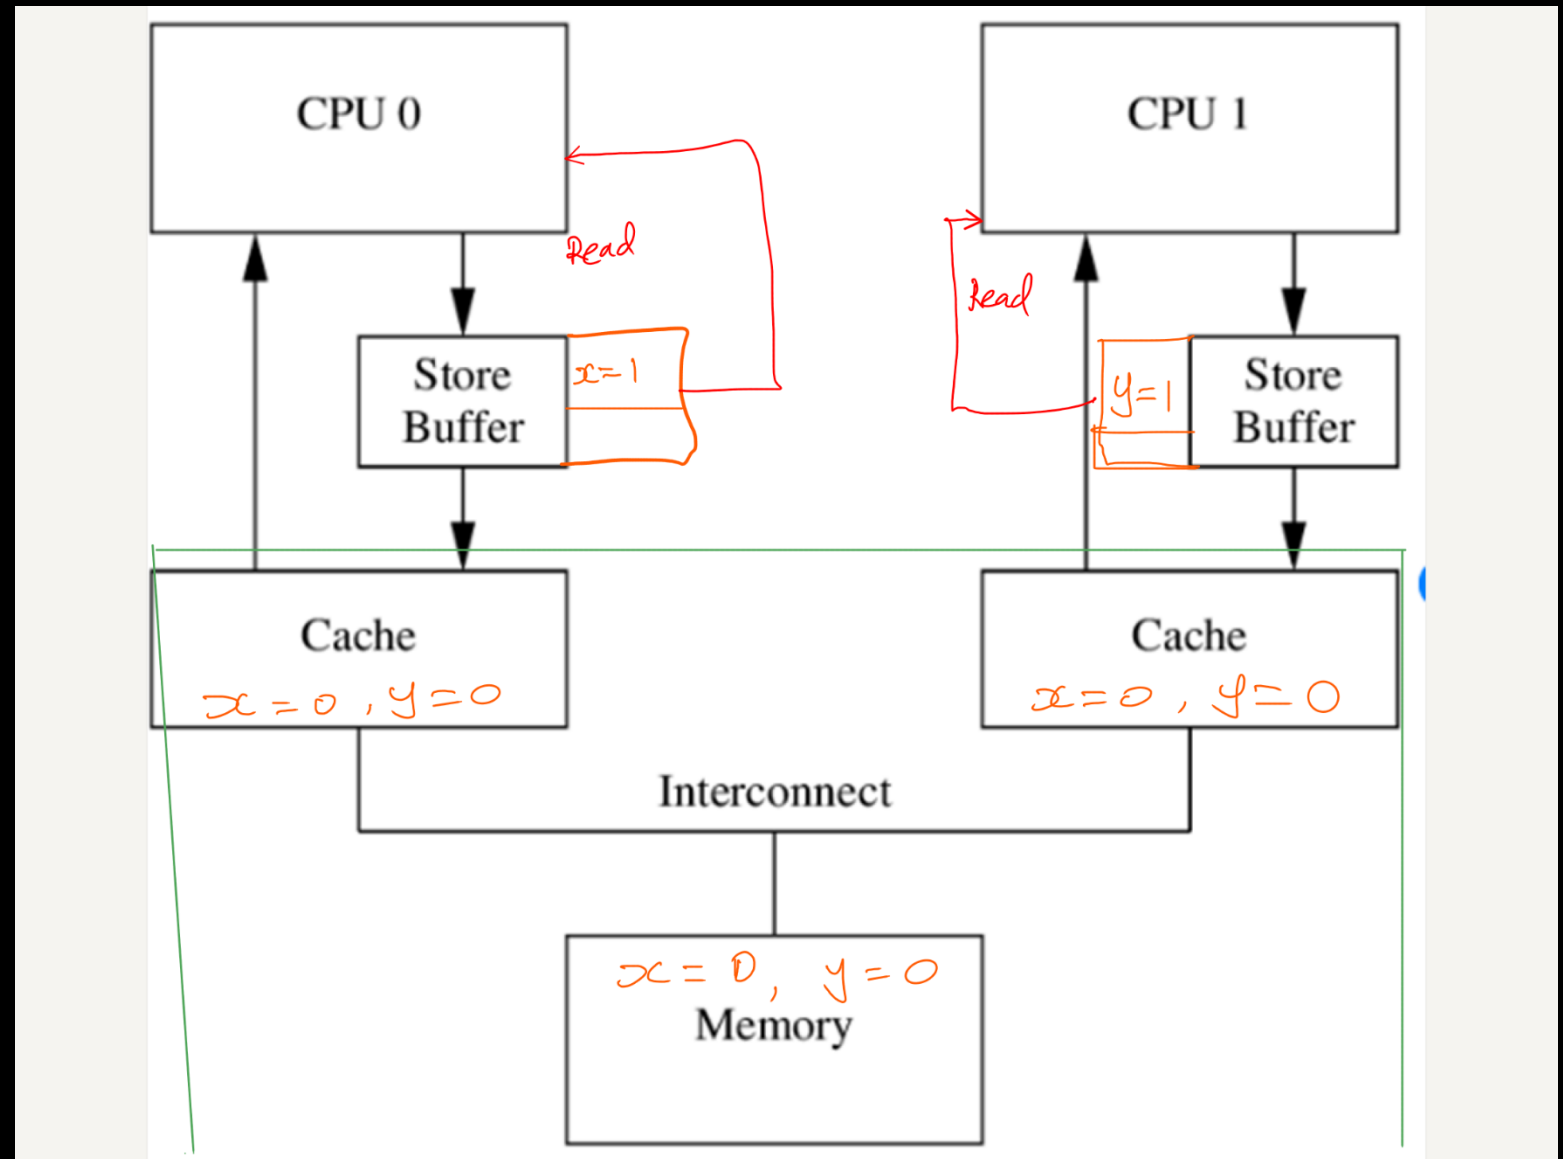
Store Buffer

- Unlike Cache this affects correctness
- Non-Coherent part of x86 memory



Store Buffer

- Write is in Store buffer
- Non-Coherent part of x86 memory
- Other core cannot observe the writes




Memory Barriers

- Flush Store buffer to coherent memory
- CPU will stall until this is done
- X86 MFENCE instruction
- X86 LOCK prefix has same effect

```
8 void *core0 (void *arg)
9 {
10     x = 1;
11     asm volatile ("mfence" ::: "memory"); /
12     r1 = y;
13     return 0;
14 }
15
16 void *core1 (void *arg)
17 {
18     y = 1;
19     asm volatile ("mfence" ::: "memory"); /
20     r0 = x;
21     return 0;
22 }
```

Memory Barriers

- X86 is one of the easiest architectures to understand
- *Only Write and Read can be re-ordered*
- Write and Write don't get re-ordered or we cannot observe from software
- Read followed by Write is not Reordered
- Read followed by Read is also not reordered
- Simple model on right is enough



Moving up Stack to JAVA

- Flush Store buffer to coherent memory
- CPU will stall until this is done
- X86 MFENCE instruction
- X86 LOCK prefix has same effect

```
29      static volatile long dummy;  
30      static void thread0()  
31      {  
32          thread0Val = 1;  
33          UNSAFE.fullFence();  
34          r1 = thread1Val;  
35  
36      }  
37
```

